

STARLIGHT

Starbucks Treatment And Location Inference for Gentrification & Housing Trajectory

Thesis

The widely-cited claim that homes near Starbucks appreciate 96% versus the 65% national average conflates two distinct effects: **(1)** Starbucks's skill at selecting neighborhoods already on an appreciation trajectory, and **(2)** the causal effect of a new Starbucks on local home prices. This project decomposes the headline number into its selection and causal components using Atlanta as a testbed, builds a doubly-robust difference-in-differences estimator that handles both confounders, and deploys the resulting models as a public API that, given any Atlanta address, returns expected appreciation properly accounting for the probability of a Starbucks opening and the genuine causal impact if one does.

Why Atlanta

- **BeltLine natural experiment.** The 22-mile trail loop opened in staggered segments starting around 2008. Cleanest urban-gentrification natural experiment in any US city, and it coincides with Starbucks's recent expansion.
- **Interpretable heterogeneity.** East-west divide, intown vs OTP, and varied gentrification trajectories give findings a clear narrative.
- **Good open data.** Fulton County, DeKalb County, Atlanta Open Data, MARTA, and the BeltLine all publish usable datasets.
- **Dramatic Starbucks variation.** Long-standing stores in Buckhead and Midtown; recent expansion into West End, Reynoldstown, Edgewood, and East Atlanta Village provides high treatment-timing variance.

High-Level Plan: Four Phases

Phase	Weeks	Goal
1. Data foundation	1–2	Clean tract × month panel with prices, treatment status, and covariates.
2. Site-selection model	3–4	Calibrated classifier predicting P(Starbucks opens) with SHAP explanations.

3. Causal model	5–6	Doubly-robust DiD producing naive, controlled, and DR treatment effects.
4. Deployment	7–8	FastAPI live, simple map frontend, README with the narrative.

Phase 1 — Data Foundation

Goal: a clean panel dataset of (tract × month) with home prices, Starbucks treatment status, and all covariates ready to model.

Geographic unit

Use **census tracts** as the primary panel unit. Small enough that a *Starbucks opened here* is meaningfully local, big enough for stable monthly price estimates, matches ACS natively, and is the standard unit for staggered DiD in urban econ. The site-selection model in Phase 2 uses a finer unit (H3 hexagons at resolution 9, ~150m across).

Datasets

Starbucks locations and opening dates (the hard one)

- **Current locations:** Starbucks store-locator JSON endpoint. ~150 stores in metro Atlanta.
- **Opening dates** (not published — triangulate from three sources):
 - **Wayback Machine + store locator:** scrape snapshots of starbucks.com/store-locator at 6-month intervals from 2010–2026; diff snapshots to find new store numbers.
 - **Fulton/DeKalb business license records:** filter for Starbucks filings.
 - **Google Places first review date:** lower bound on opening — first review is usually within 30 days.
- **Validation:** cross-reference all three sources; manually verify ~10% by Google News.

Output: `starbucks_openings.csv` with `store_number`, `lat`, `lon`, `tract_geoid`, `opening_year_month`, `source`, `confidence`.

Home prices

- **Primary panel:** Zillow ZHVI at census-tract level, monthly, 2000–present. Free CSV.
- **Validation / robustness:** Fulton and DeKalb tax assessor parcel-level transaction records.

Demographics and time-varying controls

- **ACS 5-year estimates** at tract level, 2010–2024: median income, education, race/ethnicity, age, household size, % renter, median rent.
- **LEHD LODES:** tract-level daytime population and worker inflow/outflow — critical for the site-selection model.
- **BLS QCEW:** county-level employment by industry as a metro-wide time control.

Built environment

- **OpenStreetMap:** coffee shops, restaurants, gyms, banks, grocery, schools, parks (Overpass API or Geofabrik dumps).

- **BeltLine trail openings** by segment, with dates.
- **MARTA** stations and lines (public GTFS).
- **Atlanta zoning shapefiles** for filtering eligible Starbucks sites.

Other anchor retail (confounders)

Whole Foods, Trader Joe's, Sweetgreen, Chipotle, Target — current locations as covariates; opening dates only for brands you specifically control as time-varying confounders.

Panel structure

```
tract_geoid | year_month | zhvi | treated | months_since_treatment
            | acs_median_income | acs_pct_college | daytime_pop
            | n_coffee_shops_quarter_mi | distance_to_beltline | ...
```

One row per (tract, month). ~ 500 tracts $\times \sim 150$ months $\approx$ **75K rows**. Manageable in pandas. `treated = 1` for tract-months at or after the first Starbucks opened in that tract; never-treated tracts are pure controls. Staggered-DiD methods handle this structure natively.

Phase 1 deliverable

A reproducible pipeline (Python scripts + Makefile, or Prefect/Dagster) that builds `panel.parquet` from raw sources, plus a notebook with basic EDA: tracts treated, treatment-timing distribution, parallel-trends plot of raw ZHVI.

Phase 2 — Site Selection Model

Case-control sampling

- **Positives:** every Starbucks opening location (~150 from Phase 1).
- **Negatives:** 5–10× as many comparable, retail-eligible parcels (commercially zoned, min parcel size, accessible).
- **Temporal alignment:** for each positive opened in year T, use features as of T–1; for negatives, sample random T and use T–1 features. Prevents post-treatment leakage.

Features

Built from the panel, sampled at H3 res-9 around the candidate site:

- **Catchment demographics** (ACS within 0.25, 0.5, 1.0 mi rings): income, education, age, race, household type
- **Daytime population** (LEHD): workers per acre within 0.25 mi
- **Commuter flow** (LEHD): inflow / outflow ratio
- **Co-tenancy** (OSM + retail): count of each anchor brand within 0.25 mi
- **Competitor density:** existing coffee shops within 0.1 and 0.25 mi
- **Cannibalization:** distance to nearest existing Starbucks
- **Accessibility:** nearest road class, distance to nearest signalized intersection, MARTA proximity, BeltLine proximity
- **Visibility proxies:** corner lot (boolean), parcel frontage
- **Real estate:** parcel size, building sqft (county assessor)

The model

LightGBM binary classifier, hyperparameters via Optuna with **grouped CV** (group by tract to prevent leakage of Starbucks's clustering pattern). Calibrate probabilities with isotonic regression on a held-out set so the output is a real propensity score, not a ranking. SHAP for explanations.

Validation: PR-AUC (class-imbalanced), Brier score (calibration), and a **temporal holdout** — train on openings before 2022, test on 2022–2025. Report: *of my top 100 predicted opportunity sites in 2022, X% became actual Starbucks by 2025.*

Phase 2 deliverable

Trained model artifact `site_selection_v1.lgb` + notebook with SHAP plots + a `score_location(lat, lon)` function returning a calibrated $P(\text{Starbucks opens})$ + temporal-holdout numbers.

Phase 3 — Causal Model

This is where the rigor lives. Three nested estimates, each strictly more credible than the last.

Estimate 1: Naive DiD

Standard two-way fixed effects on the panel:

```
log(zhvi) ~ treated + tract_FE + month_FE
```

Report the coefficient on `treated` as the naive effect. This will be inflated — *that's the point*. It is the baseline you are going to beat.

Estimate 2: Staggered DiD (Callaway–Sant'Anna)

Use the `differences` Python package (or call R's `did` via `rpy2`). CS handles staggered treatment timing without the bias that pure two-way FE suffers when effects vary.

Report:

- Overall ATT (average treatment effect on the treated)
- Event-study plot: effect by months-since-treatment; pre-period = parallel-trends test
- Heterogeneous effects by cohort (tracts treated in 2018 vs 2022, etc.)

Estimate 3: Doubly-Robust DiD

This is the headline result. Combine:

- The site-selection model (Phase 2) as the **propensity score** for treatment
- A regression model for the outcome (log price) given covariates

The DR estimator (Sant'Anna–Zhao 2020) reweights treated and control units by the inverse of the site-selection propensity, then runs DiD on the reweighted sample. **Key property:** if *either* the propensity model or the outcome model is correct, the estimate is consistent.

Headline finding

```
selection_bias_share = (estimate_2 - estimate_3) / estimate_2
```

If this is, say, 0.6, the README says: *60% of the apparent Starbucks-appreciation premium reflects Starbucks's site-selection skill, not a causal effect of new stores on prices.*

Phase 3 deliverable

Notebook with all three estimates, event-study plots, parallel-trends robustness checks, the placebo test (announced-but-not-opened locations), and a plain-language interpretation. Saved estimates feed the FastAPI `/predict_impact` endpoint.

Phase 4 — Deployment

Stack

- **FastAPI** for the API
- **Pydantic** for request/response schemas
- **Uvicorn** for the server
- **Render** or **Fly.io** for hosting (free tiers, both fine for FastAPI)
- **Leaflet + single-page HTML/JS** served as a static file from the same FastAPI app
- **OpenStreetMap** or **Mapbox tiles** for the basemap
- Skip Docker for V1 unless already comfortable — `requirements.txt` + Render config is enough.

Endpoints

```
POST /score_location      -> site-selection score + SHAP top features
POST /predict_impact      -> stacked: P(opens) x causal effect
GET /opportunity_zones     -> top-N hexes ranked by expected impact
GET /tract_summary/{geoid} -> all data for one tract (map detail view)
```

Example /predict_impact response

```
{
  "starbucks_viability_score": 0.78,
  "p_starbucks_opens_within_5yr": 0.34,
  "predicted_appreciation_if_opens": {
    "naive_did_estimate": 8.4,
    "doubly_robust_estimate": 3.1,
    "selection_bias_share": 0.63
  },
  "expected_appreciation": 1.05,
  "interpretation": "Moderate viability, but selection bias means
                    the headline number overstates the true causal
                    effect by ~63%."
}
```

$\text{expected_appreciation} = p_opens \times \text{causal_effect}$ — integrating over whether Starbucks actually opens. A quantity neither sub-model produces alone.

Map demo

Single-page Leaflet map of Atlanta. Tracts colored by expected appreciation. Click a tract → side panel shows the decomposition. Toggle existing-Starbucks layer. **This is your portfolio screenshot.**

Phase 4 deliverable

Live URL · README with the narrative structure · reproducible repo · 5-minute Loom walkthrough (optional, high ROI).

README Narrative Structure

The story flow that hiring managers will remember:

1. **The hook.** Everyone has seen the 96% vs 65% appreciation stat. What's actually going on?
2. **The naive answer.** Basic correlation. Confirms the stat.
3. **First refinement.** DiD around store openings. Effect shrinks but stays large.
4. **The selection problem.** Starbucks does not open at random. Train a model on their historical site selection to learn *how* they pick.
5. **Findings from the site model.** What Starbucks actually weights — daytime population, co-tenancy with X, commuter flow — with SHAP plots.
6. **The doubly-robust answer.** Combining the propensity with DiD gives a much smaller causal effect. The headline overstates by ~Xx.
7. **The product.** Deployed as a tool that, given any address, returns expected appreciation properly accounting for both whether Starbucks would open and how much it would actually move prices.
8. **The bonus.** Opportunity-zone visualization showing high-impact under-served locations.

Repo Structure

```
starlight/
■■■ README.md
■■■ pyproject.toml           # or requirements.txt
■■■ Makefile                 # for reproducibility
■■■ data/
■   ■■■ raw/                 # gitignored, populated by scripts
■   ■■■ interim/
■   ■■■ processed/
■       ■■■ panel.parquet
■■■ src/
■   ■■■ data/
■       ■■■ starbucks_scraper.py
■       ■■■ wayback_scraper.py
■       ■■■ census_pull.py
■       ■■■ osm_pull.py
■       ■■■ build_panel.py
■   ■■■ models/
■       ■■■ site_selection.py
■       ■■■ causal_did.py
■       ■■■ stacked.py
■   ■■■ api/
■       ■■■ main.py          # FastAPI app
■       ■■■ schemas.py       # Pydantic models
■       ■■■ inference.py     # model loading + prediction logic
■   ■■■ frontend/
■       ■■■ index.html       # Leaflet map
■■■ notebooks/
■   ■■■ 01_eda.ipynb
```



```

■ ■■■ 02_site_selection.ipynb
■ ■■■ 03_causal_naive.ipynb
■ ■■■ 04_causal_did.ipynb
■ ■■■ 05_causal_doubly_robust.ipynb
■ ■■■ 06_validation.ipynb
■■■ models/                # serialized .lgb, .pkl artifacts
■■■ tests/
    ■■■ test_api.py

```

Tools / Libraries

Python 3.11+ · pandas · geopandas · h3-py · lightgbm · optuna · shap · statsmodels · **differences** (Callaway–Sant’Anna and DR-DiD in Python — the key library for Phase 3) · scikit-learn · fastapi · pydantic · uvicorn · httpx (scraping). Visualization: matplotlib · plotly. Frontend: Leaflet + vanilla JS (no React needed).

Realistic Risks and Mitigations

Risk: Wayback Machine scraping for opening dates is brittle.

Mitigation: Have the county business-license fallback ready from week 1; if Wayback is noisy, lean on that as the primary.

Risk: Parallel trends fails for many treated tracts.

Mitigation: This is a *finding*, not a failure. It means selection bias is real and is exactly why the doubly-robust estimator matters. Lean into it in the README.

Risk: DR-adjusted causal effect is close to zero or negative.

Mitigation: Also a valid finding. Headline becomes *the Starbucks effect is mostly selection* rather than *Starbucks causes X% appreciation*. Either result is publishable-quality.

Risk: Scope creep on site-selection features.

Mitigation: Set a feature-set budget in week 3 and stop adding once you hit it.

This Week

- **Day 1–2:** set up the repo. Get the Starbucks store-locator scraper working for current Atlanta locations. Confirm Census ACS API access. Pull Atlanta tract shapefiles.
- **Day 3–4:** start the Wayback Machine scraper for historical Starbucks snapshots. Run it in the background — one snapshot per quarter from 2010 to present.

- **Day 5–7:** download ZHVI tract-level data; basic EDA — how many Atlanta tracts ever got a Starbucks, what does the raw appreciation gap look like, plot on a map.

That EDA at the end of week 1 is your decision point. If Atlanta's tract-level data clearly shows the appreciation gap, proceed with confidence. If it doesn't, reconsider city choice before sinking more time.

Project plan generated for an undergraduate ML / informatics portfolio piece. Atlanta scope V1 → Houston/national replication V2.